

Kubernetes I Deep Dive

Agenda

Part I

- Why Kubernetes
- What is Kubernetes
- Installation
- kubectl
- Pod
- ReplicaSet
- Deployment

Part II

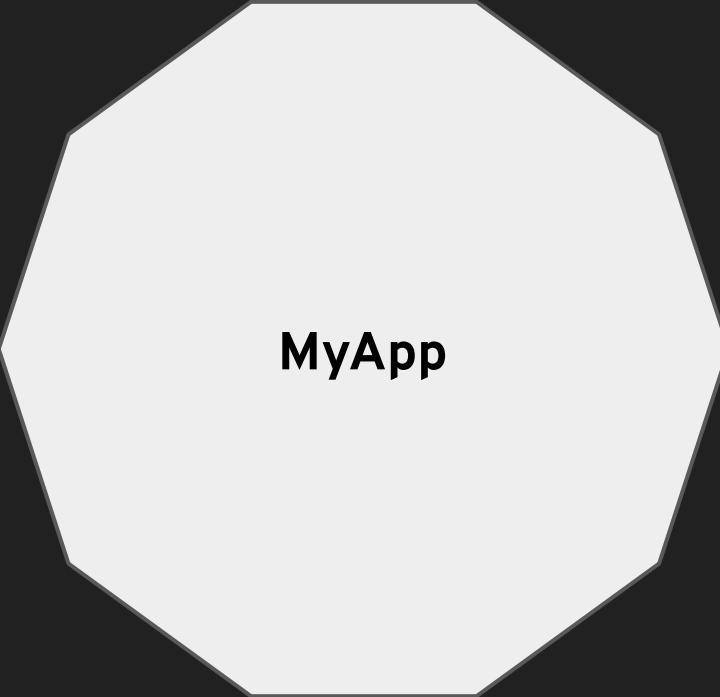
- Building Images
- Resource Limits
- Rolling Updates
- Live & Ready
- Env & ConfigMap

Part III

- Volumes
- Secrets
- Operators
- Taints & Affinity
- Jobs & CronJobs
- DaemonSet
- StatefulSets

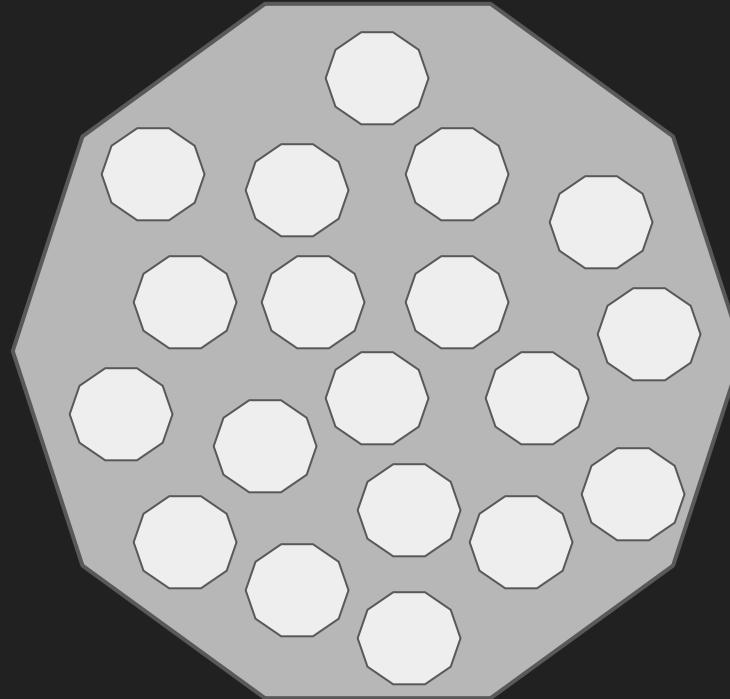
Why Kubernetes?

Monolith

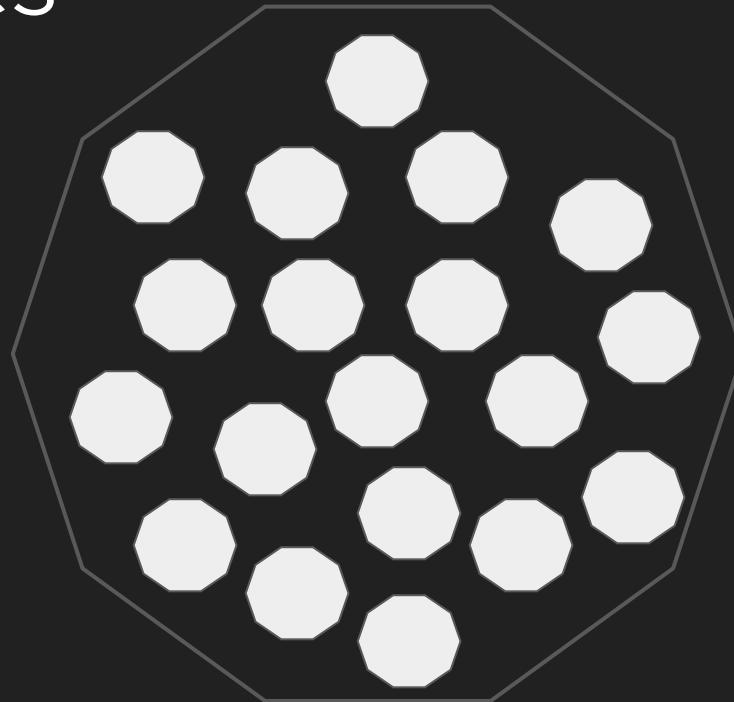


MyApp

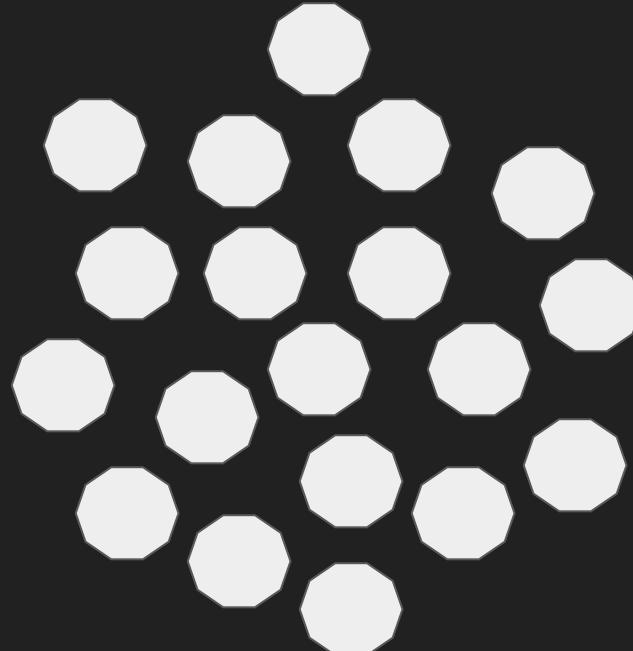
Application Modules

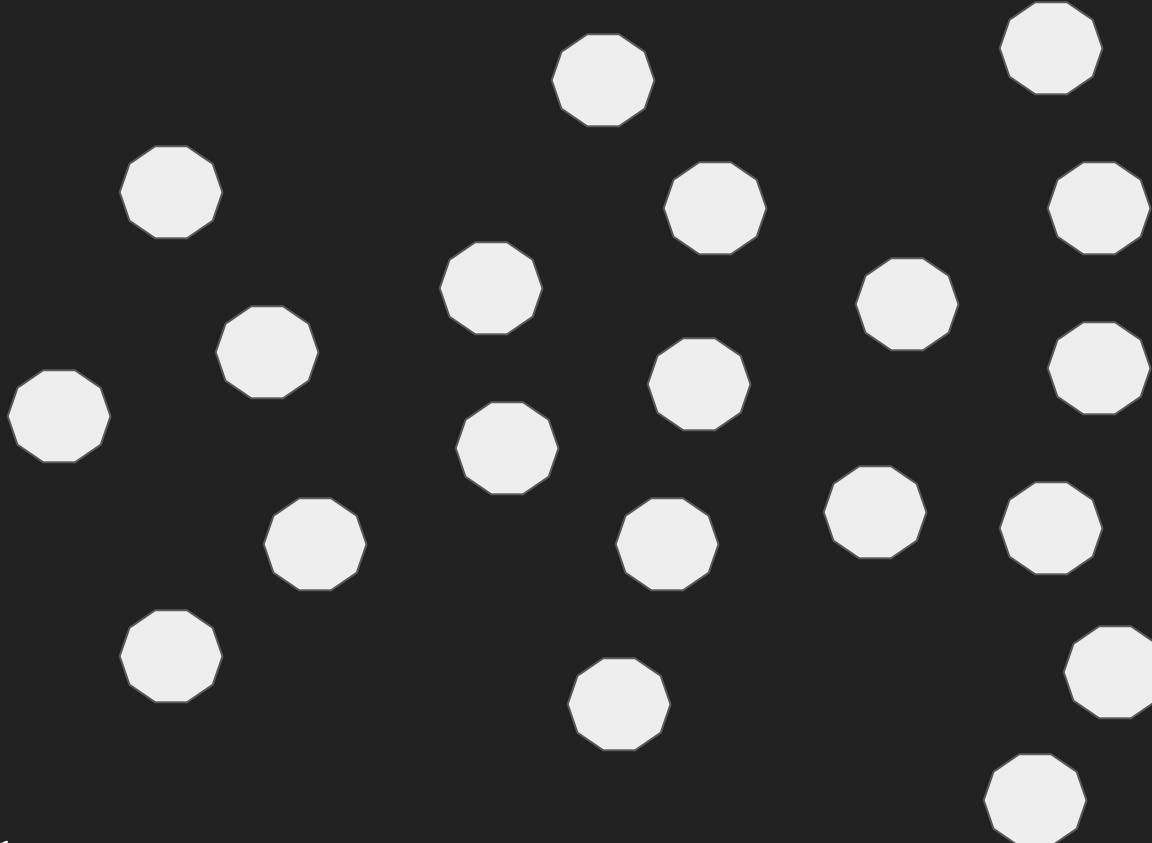


Microservices

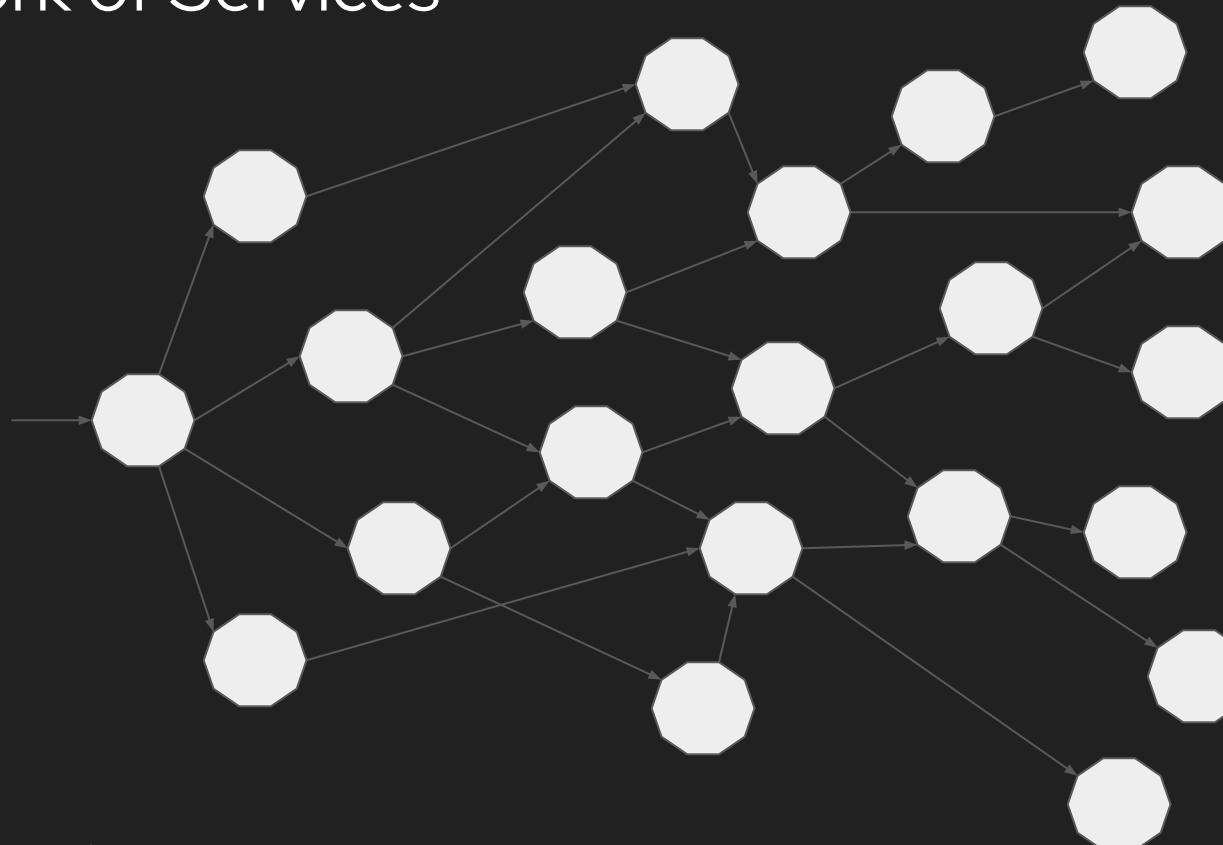


Containers

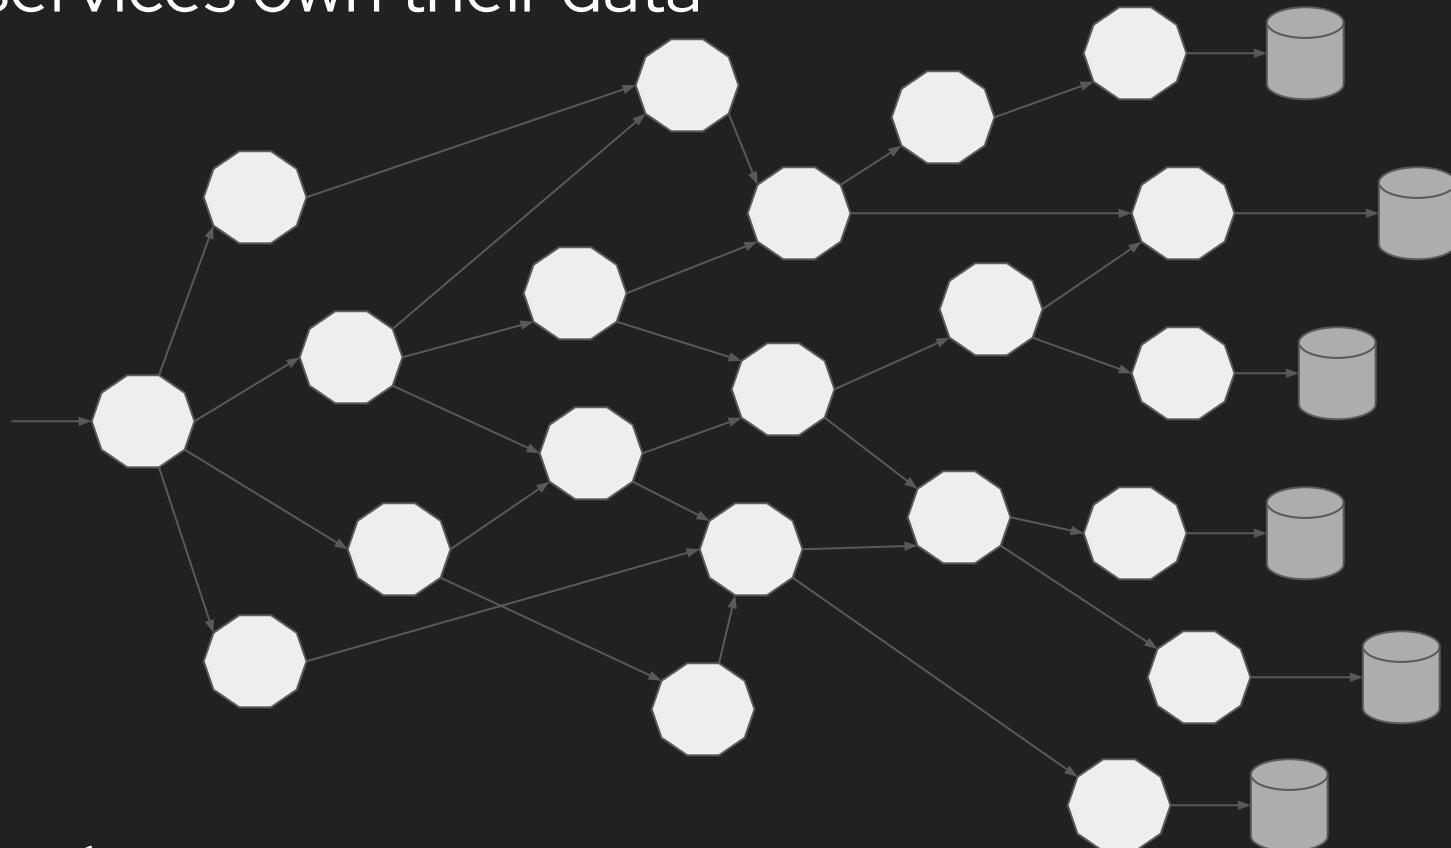




Network of Services



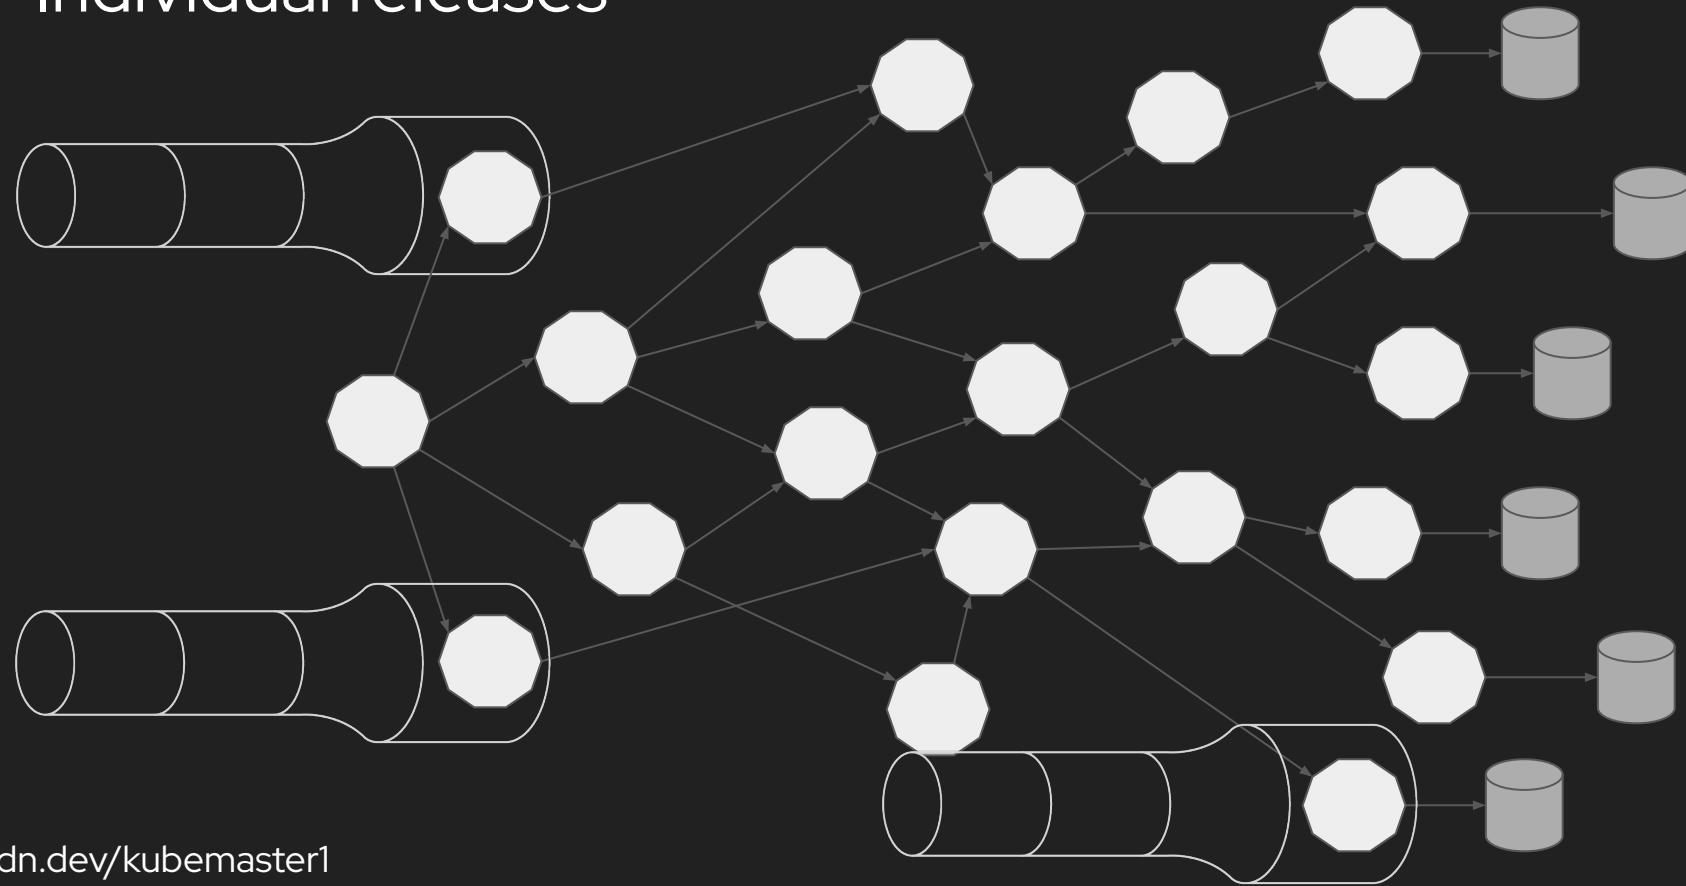
Microservices own their data



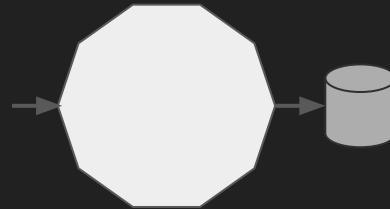
Multiple points of entry



Individual releases

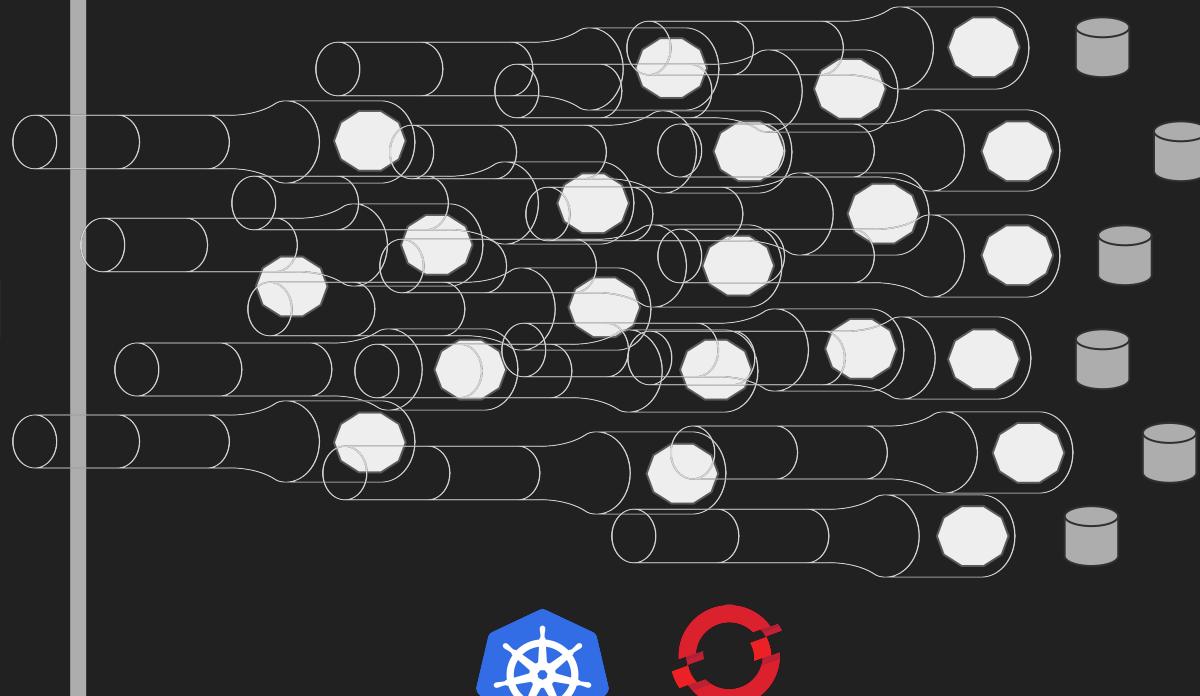


Every 3 months



Love Thy Mono

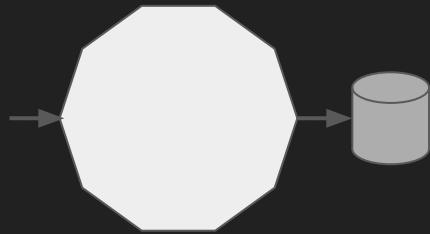
Every 3 weeks/days/hours



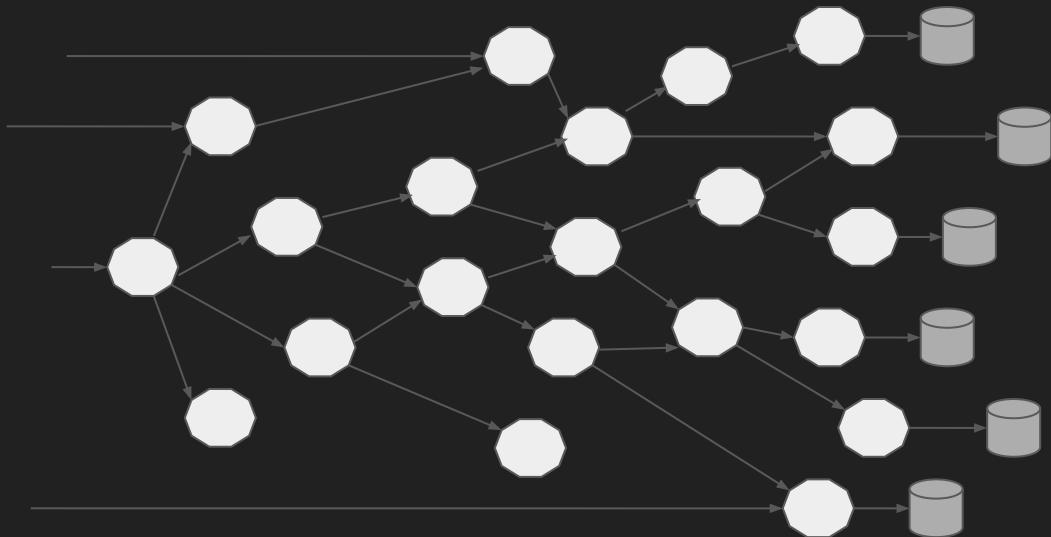
OPENSHIFT

dn.dev/kubemaster1

Old School



New School



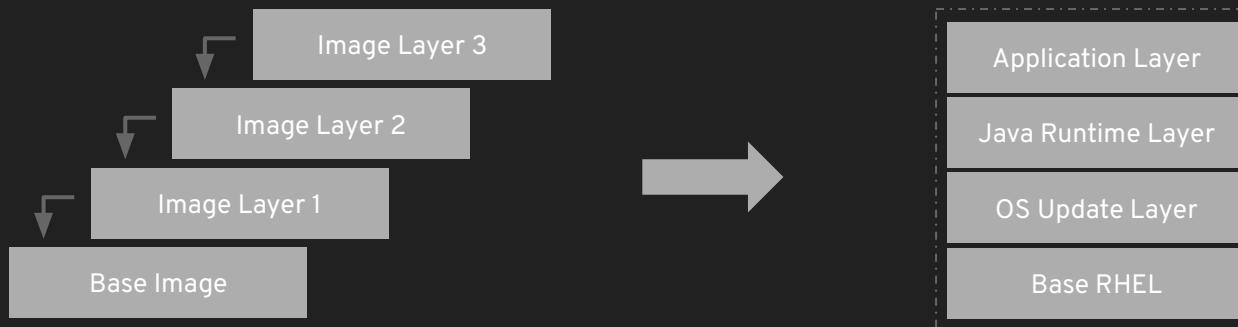
OPENSHIFT

dn.dev/kubemaster1

Containers



Container Image Layers



Container Image Layers

Example Container Image



Dockerfile: create your Container

- 1 Inherit from a base image
- 2 Parameters as environment variables
- 3 Install dependencies (tooling from base image)
- 4 Add your app as a new Layer
- 5 Expose the port your app will use
- 6 Run the app

```
FROM registry.access.redhat.com/ubi8/ubi
ENV foo=text
RUN dnf install -y java-11-openjdk
ADD my-app.jar /home/my-app.jar
EXPOSE 8080
CMD java -jar /home/my-app.jar
```



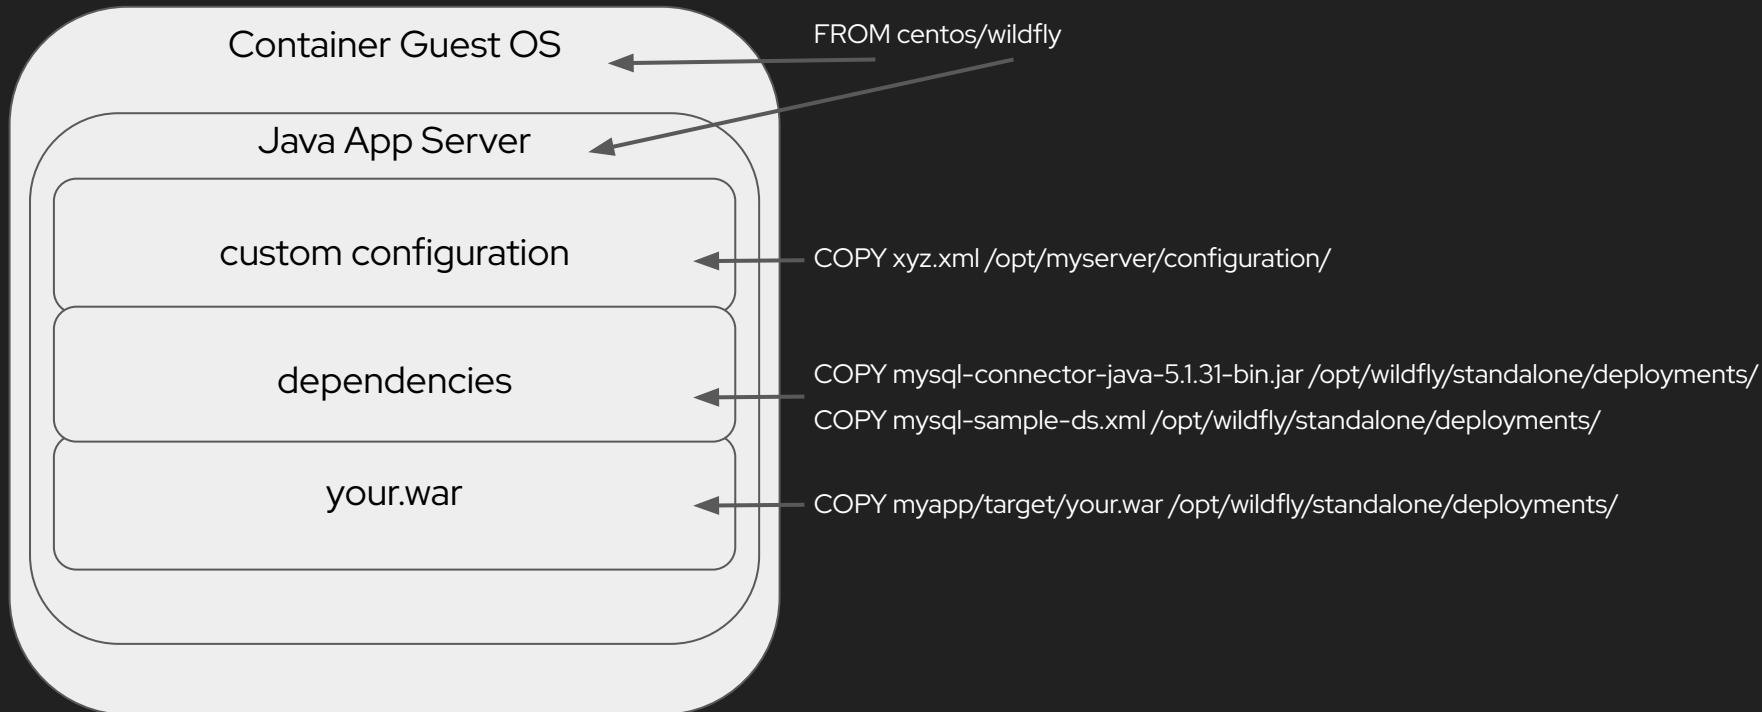
A Challenge

Have you ever had “/” vs “\” break your app? Or perhaps needed a unique version of a JDBC driver? Or had a datasource with a slightly misspelled JNDI name? Or received a patch for the JVM or app server that broke your code?

Containerize
Your
App

.war or .ear	
Custom Configuration	JDBC driver, datasource, JMS queue, users
Application Server	Weblogic 10.x.y, Tomcat 6.x.y, JBoss EAP 6.x.y
Java Virtual Machine	Java 1.6.6_45 or Java 1.7.0_67
Operating System	Linux Kernel Version & Distribution
Server Hardware	

Anatomy of a Container Image



DevOps Challenges for Multiple Containers

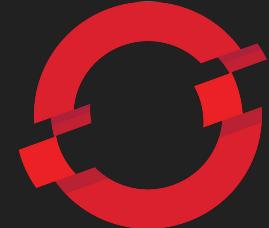
- How to scale?
- How to avoid port conflicts?
- How to manage them on multiple hosts?
- What happens if a host has trouble?
- How to keep them running?
- How to update them?
- Rebuild Container Images?



Kubernetes

Meet Kubernetes

- Greek for “Helmsman,” also the root of the word “Governor” (from latin: gubernator)
- Container orchestrator
- Supports multiple cloud and bare-metal environments
- Open source, written in Go
- Manage applications, not machines

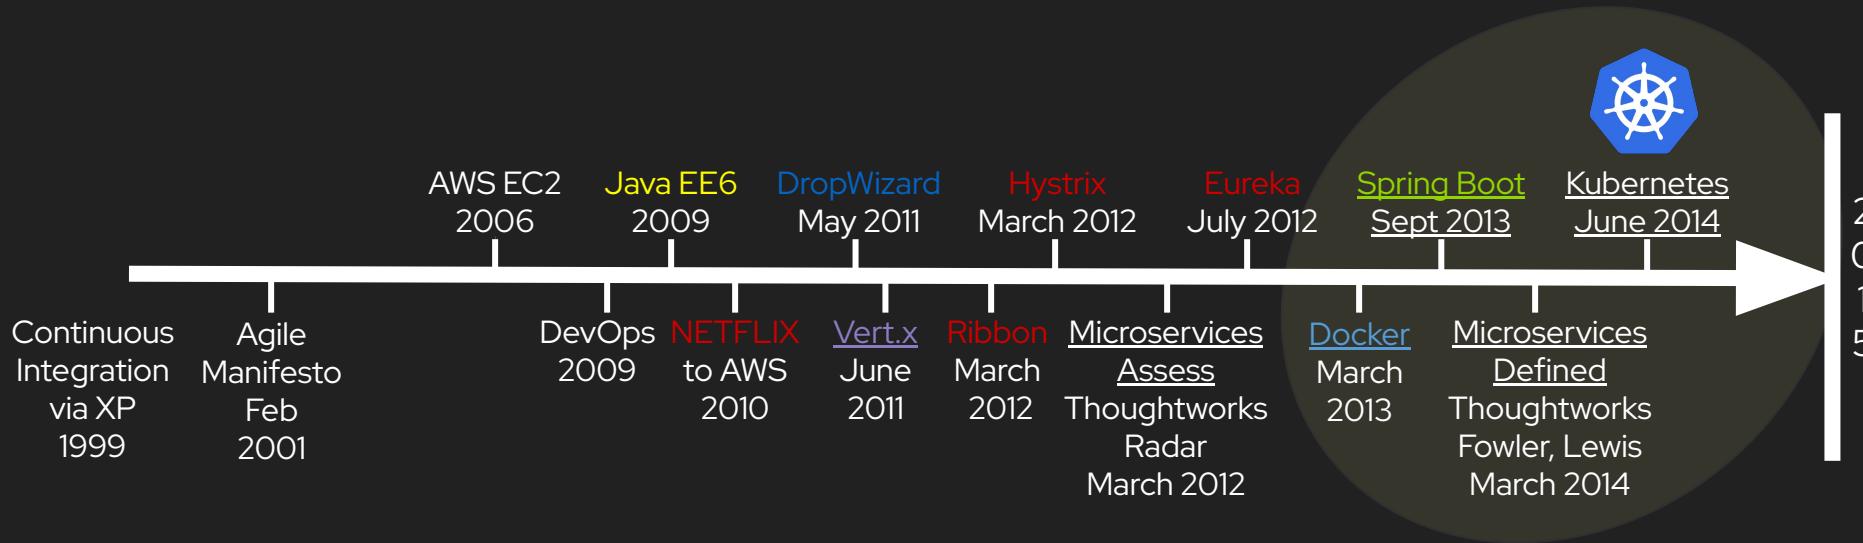


OPENSHIFT



[Learn more](#)

History of Microservices

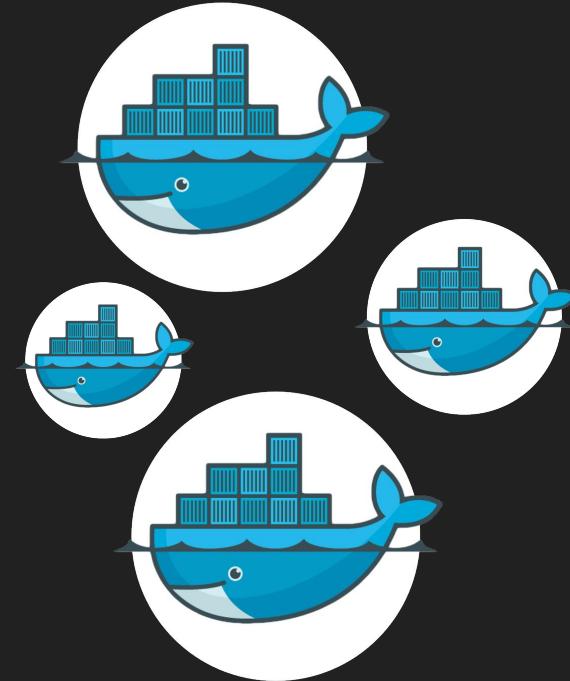


What is Kubernetes?

Pods

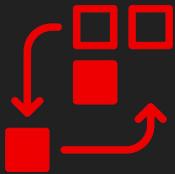
A group of whales is commonly referred to as a pod and a pod usually consists a group of whales that have bonded together either because of biological reasons or through friendships developed between two or more whales.

In many cases a typical whale pod consists of anywhere from 2 to 30 whales or more.*



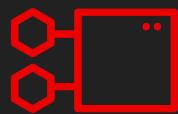
*<http://www.whalefacts.org/what-is-a-group-of-whales-called/>

Pod

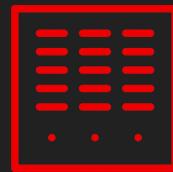
ReplicaSet/
Deployment

- ✓ 1+ containers
- ✓ Shared IP
- ✓ Shared storage (ephemeral)
- ✓ Shared resources
- ✓ Shared life cycle

Service



dn.dev/kubemaster1

Persistent
Volume

- ✓ Network
- ✓ available storage
- ✓ PVs and PVCs

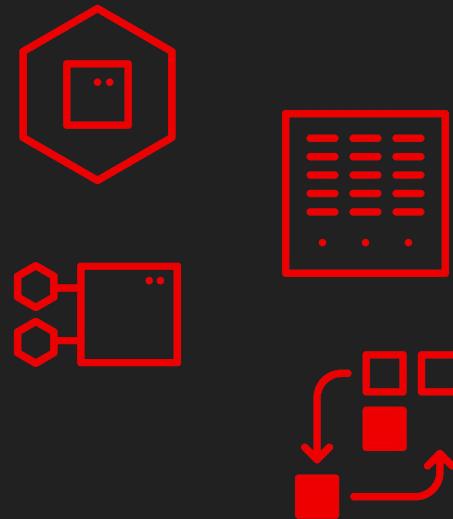
Label



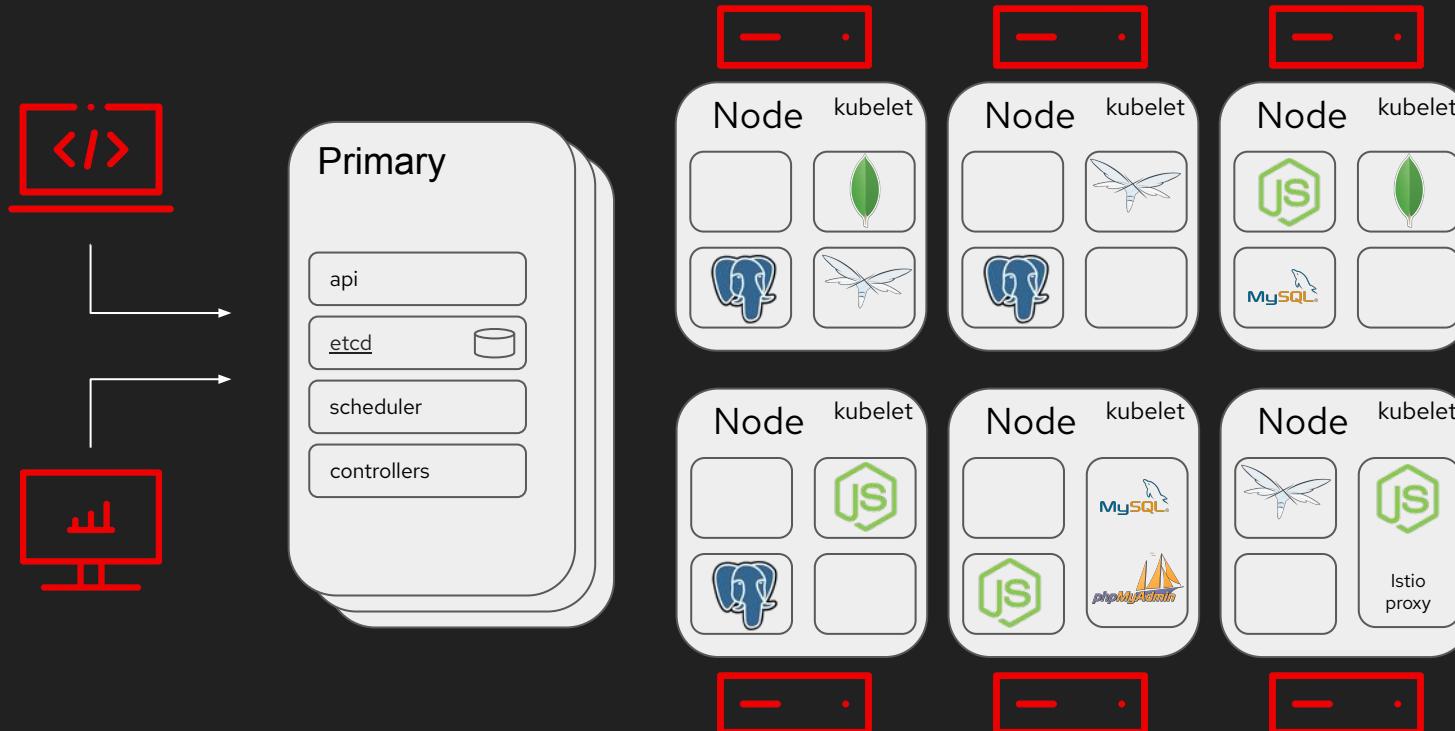
- ✓ Key/value pairs associated with Kubernetes objects (env=production)

Namespaces

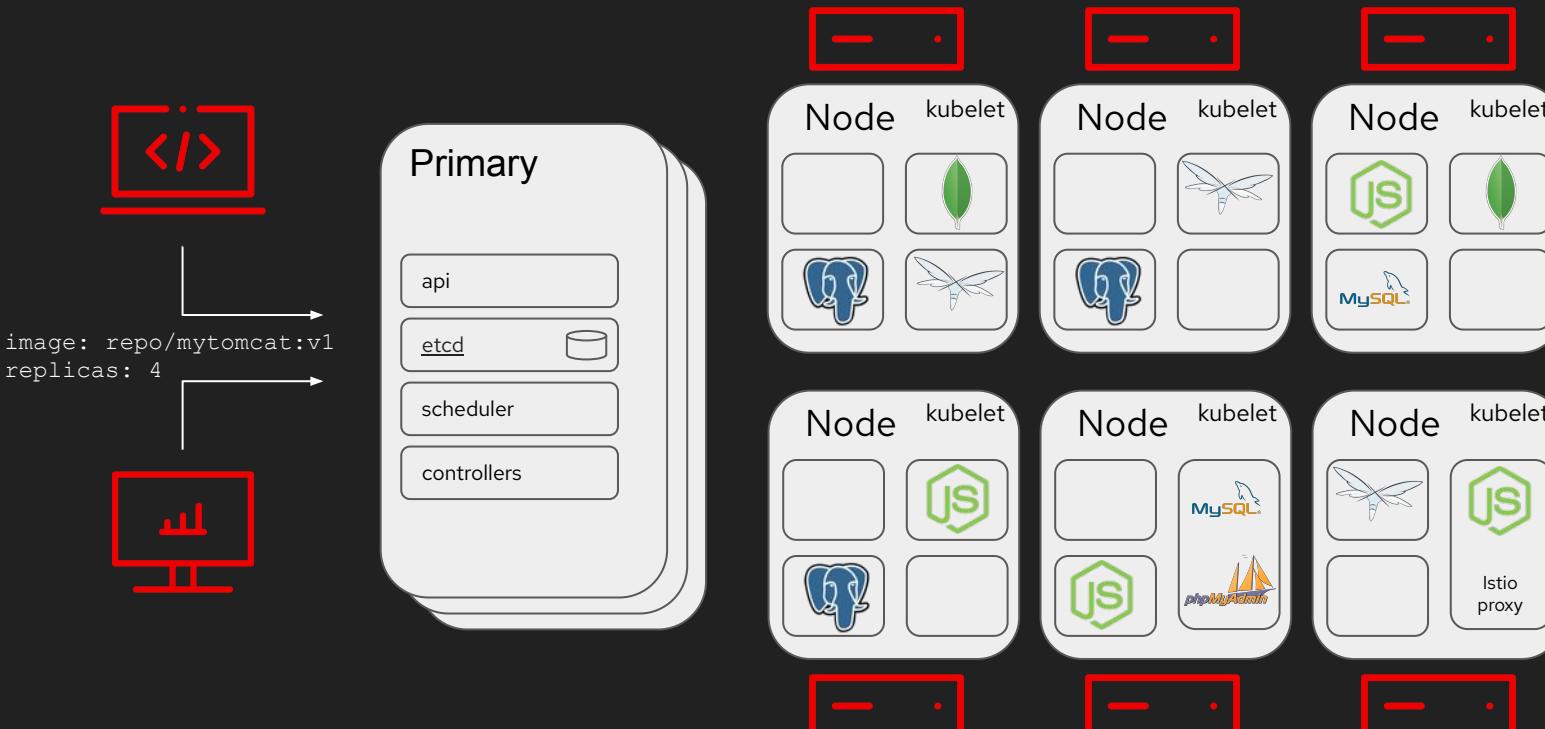
- ✓ A way to divide cluster resources between multiple users
- ✓ Support attachment of authorization and usage policies



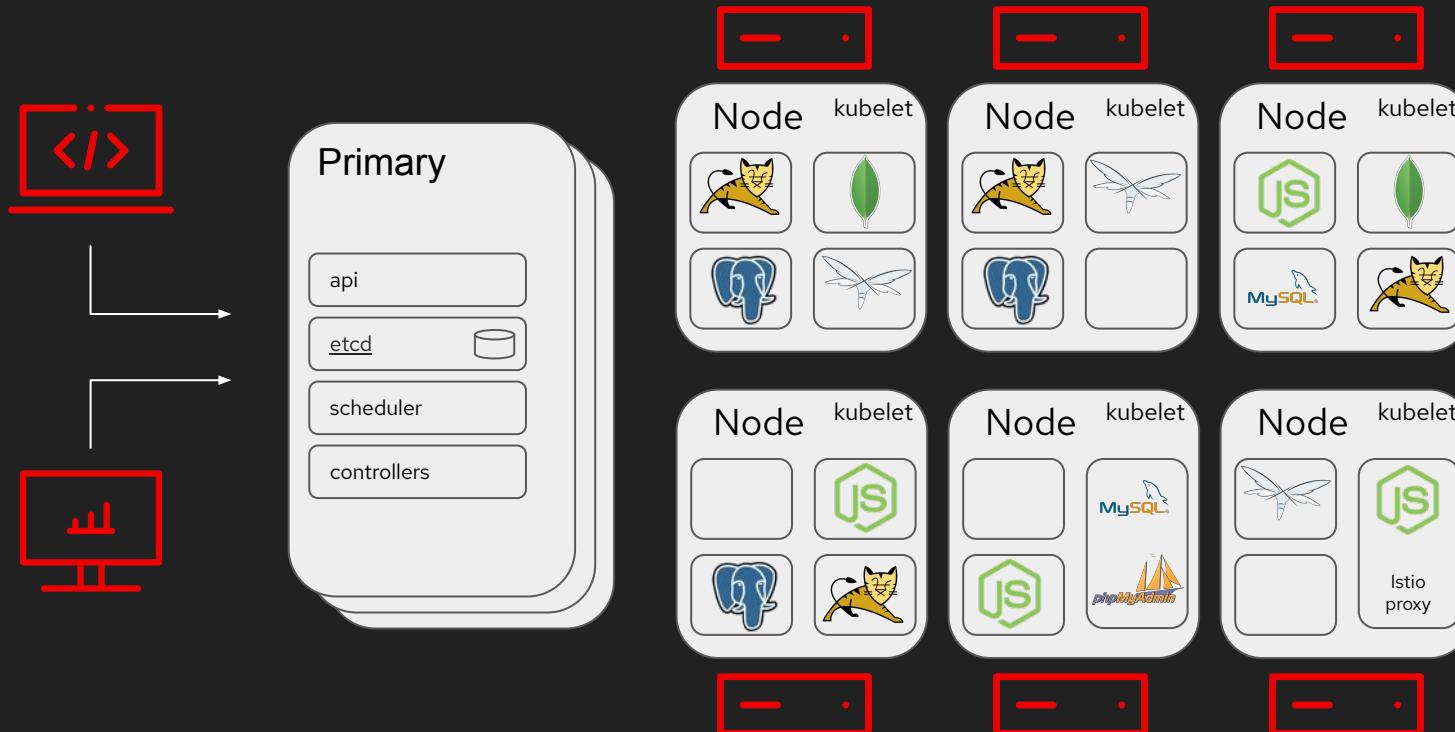
Kubernetes Cluster - Nodes



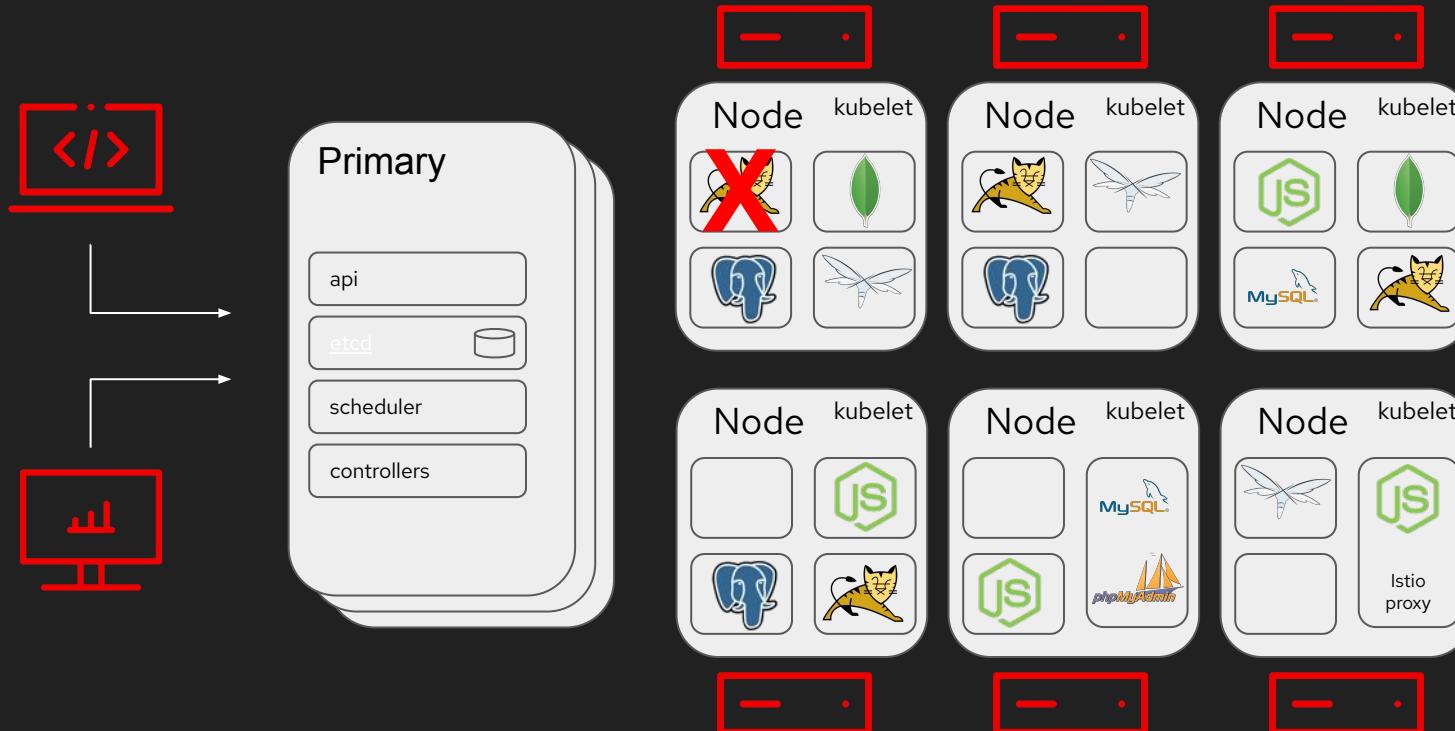
Kubernetes Cluster - Declarative API



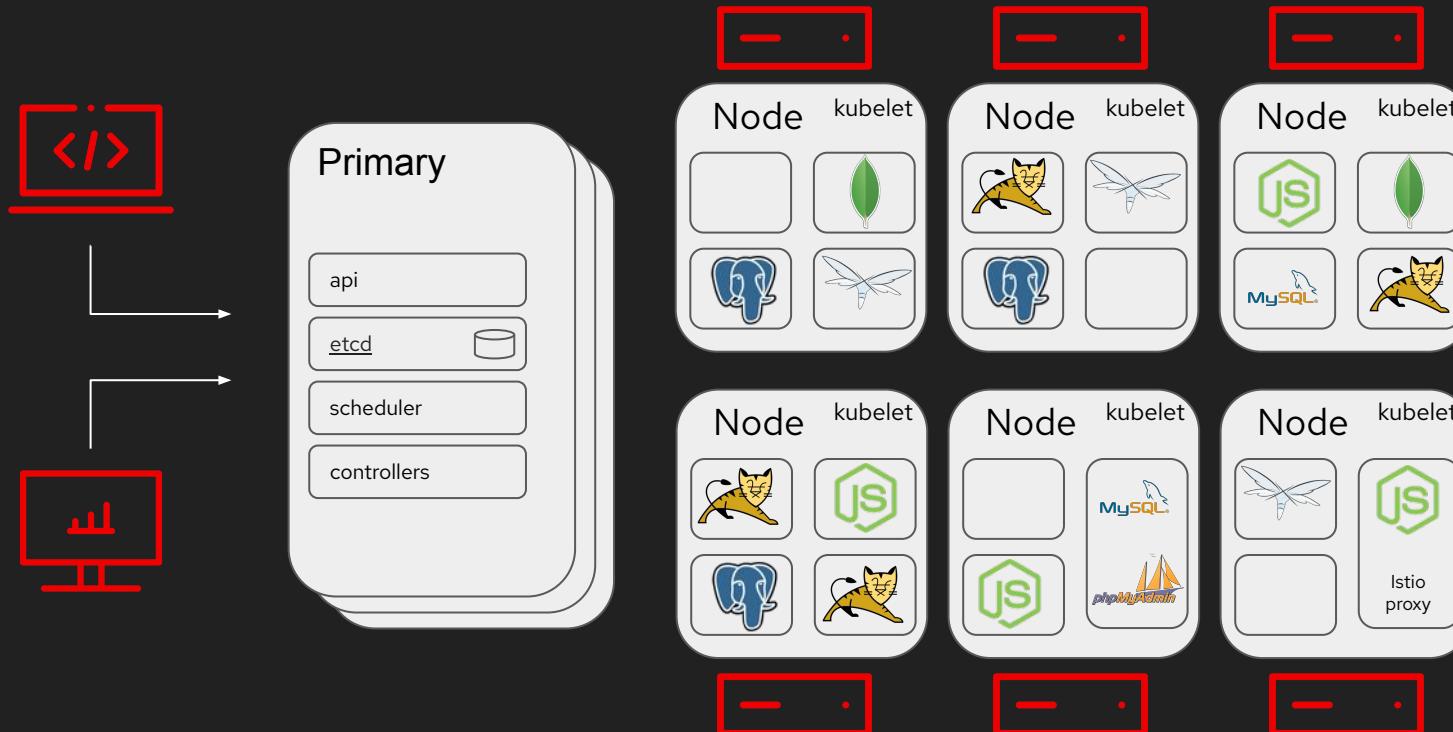
Kubernetes Cluster - 4 Tomcat Instances



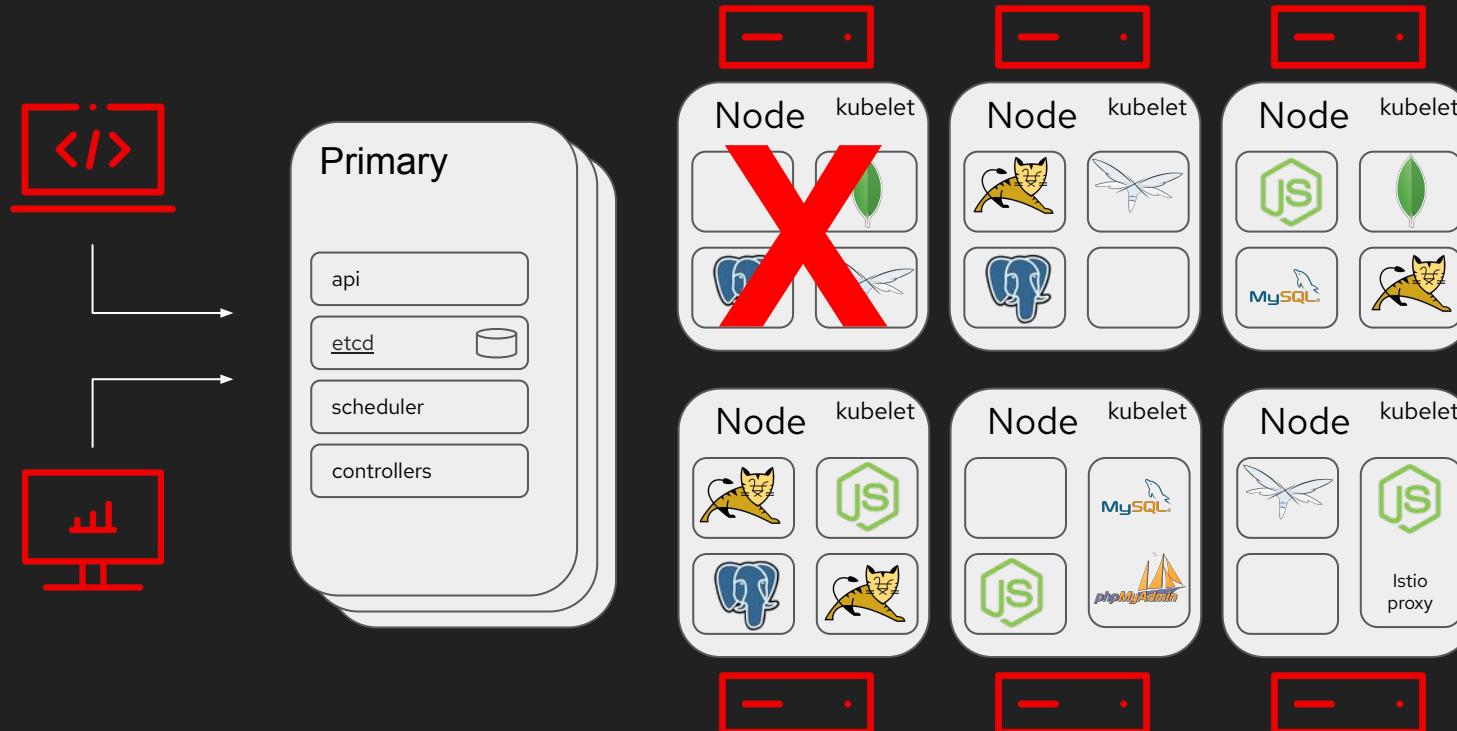
Kubernetes Cluster - Pod Failure



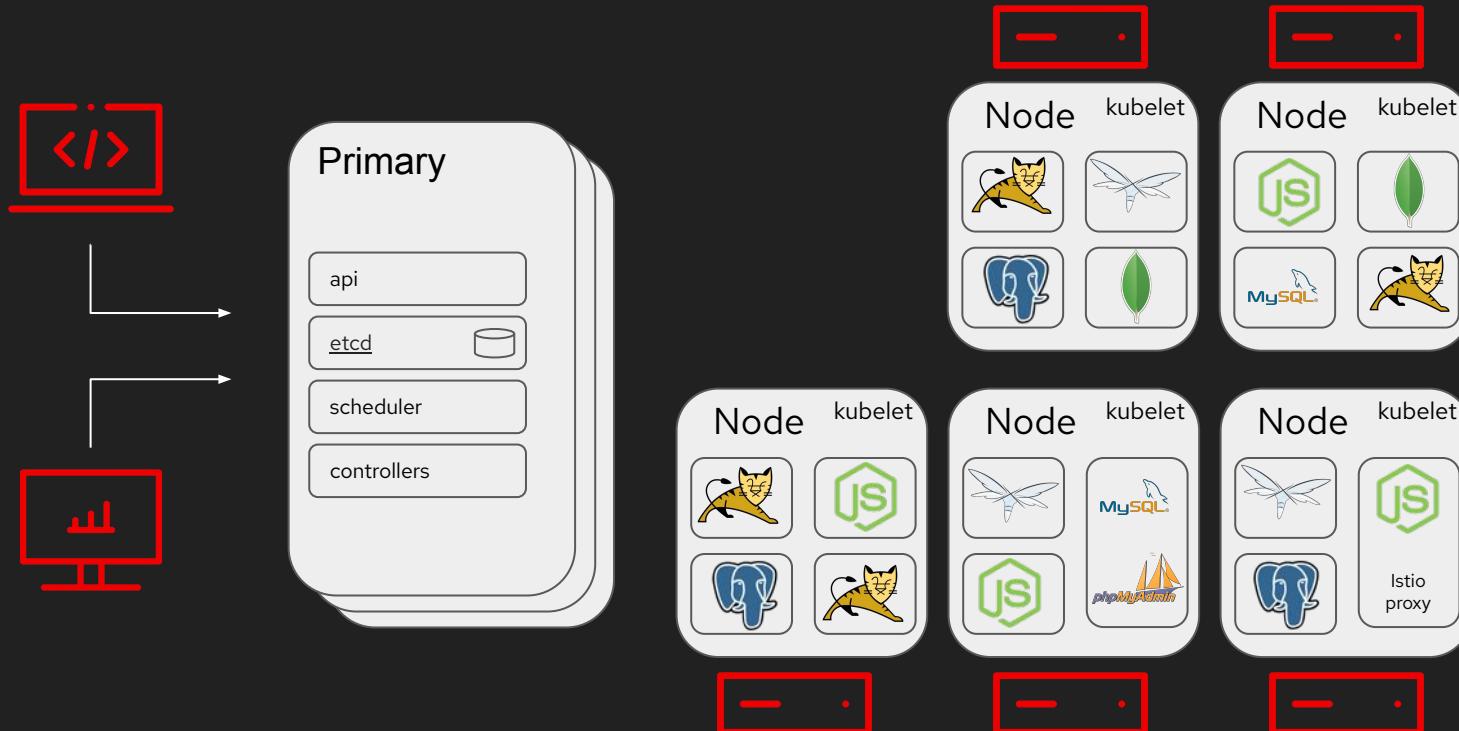
Kubernetes Cluster - Recovery



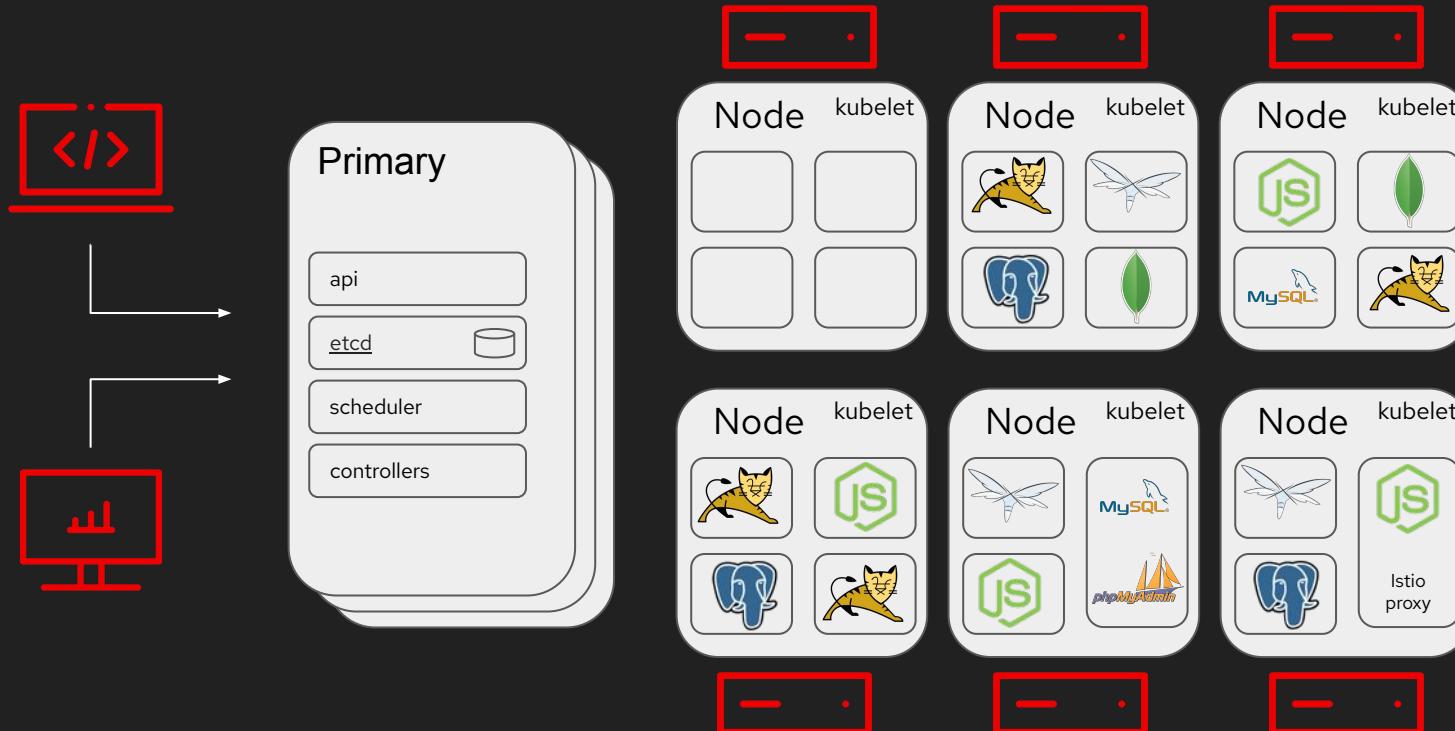
Kubernetes Cluster - Node Failure



Kubernetes Cluster - Pods Replaced



Kubernetes Cluster - New Node



Pod, ReplicaSet, Deployment

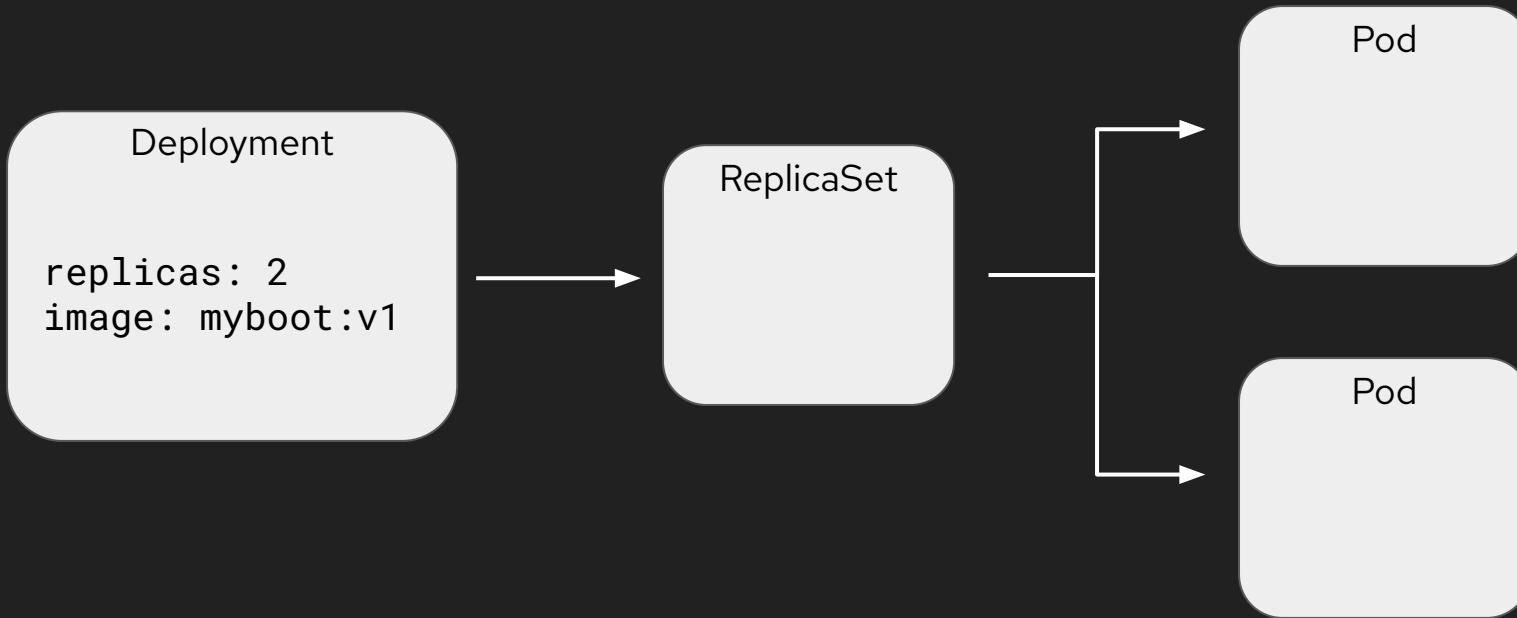
Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: myboot-demo
spec:
  containers:
  - name: myboot-demo
    image: quay.io/rhdevelopers/myboot:v4
```

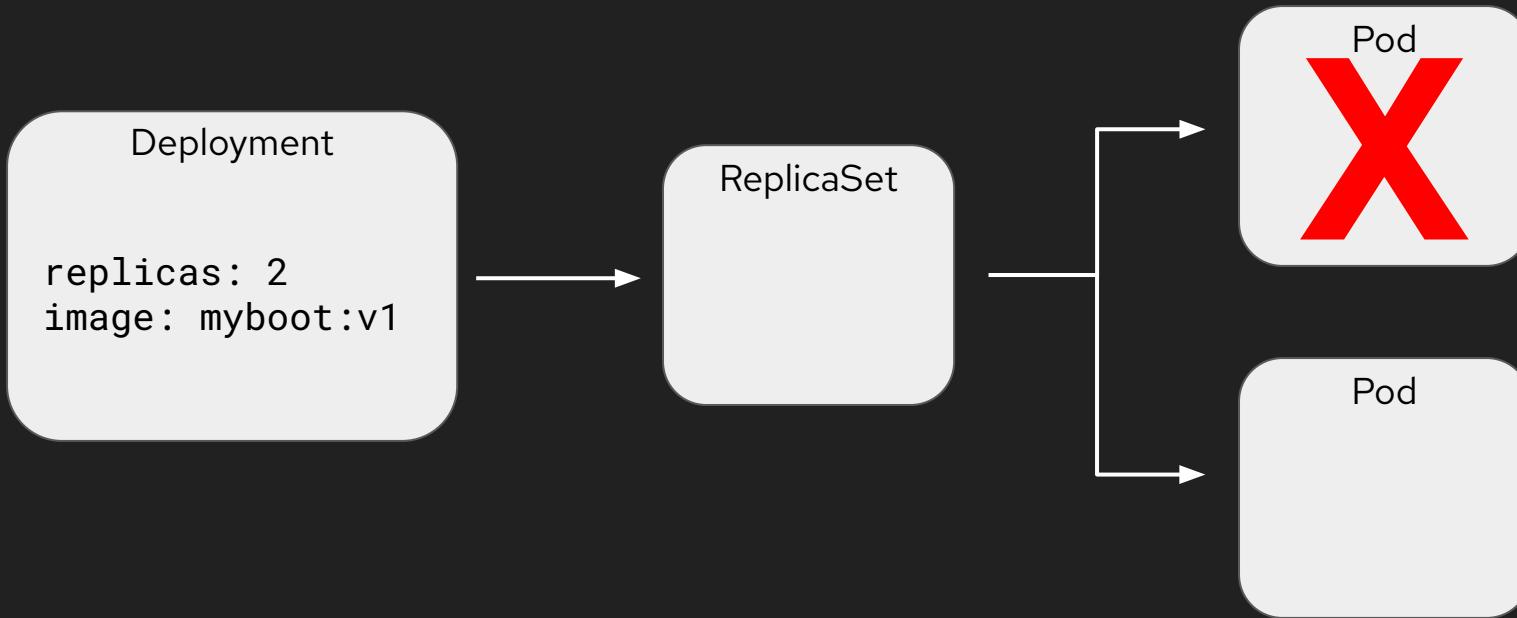
Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myboot
spec:
  replicas: 1
  selector:
    matchLabels:
      app: myboot
  template:
    spec:
      containers:
        - name: myboot
          image: quay.io/rhdevelopers/myboot:v1
          ports:
            - containerPort: 8080
```

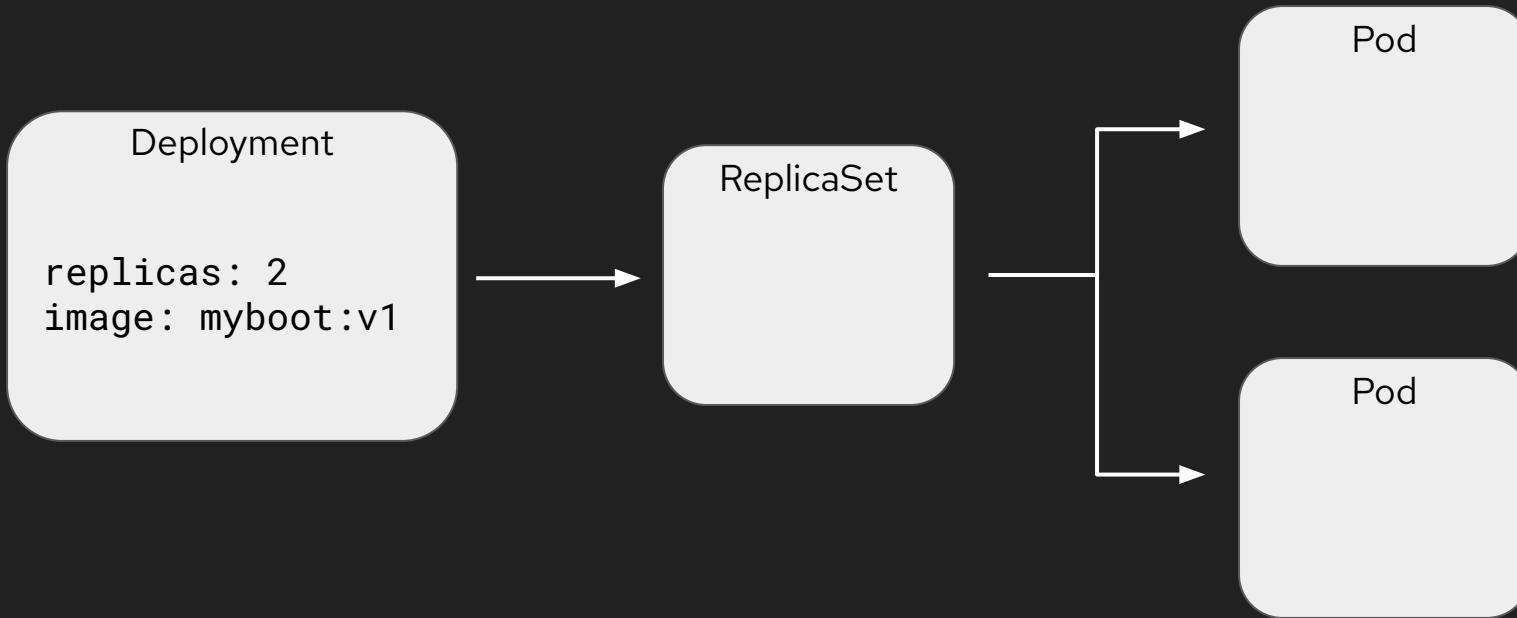
Deployments



Deployments



Deployments



Rolling Updates

“Rolling updates allow Deployments' update to take place with zero downtime by incrementally updating Pods instances with new ones.”

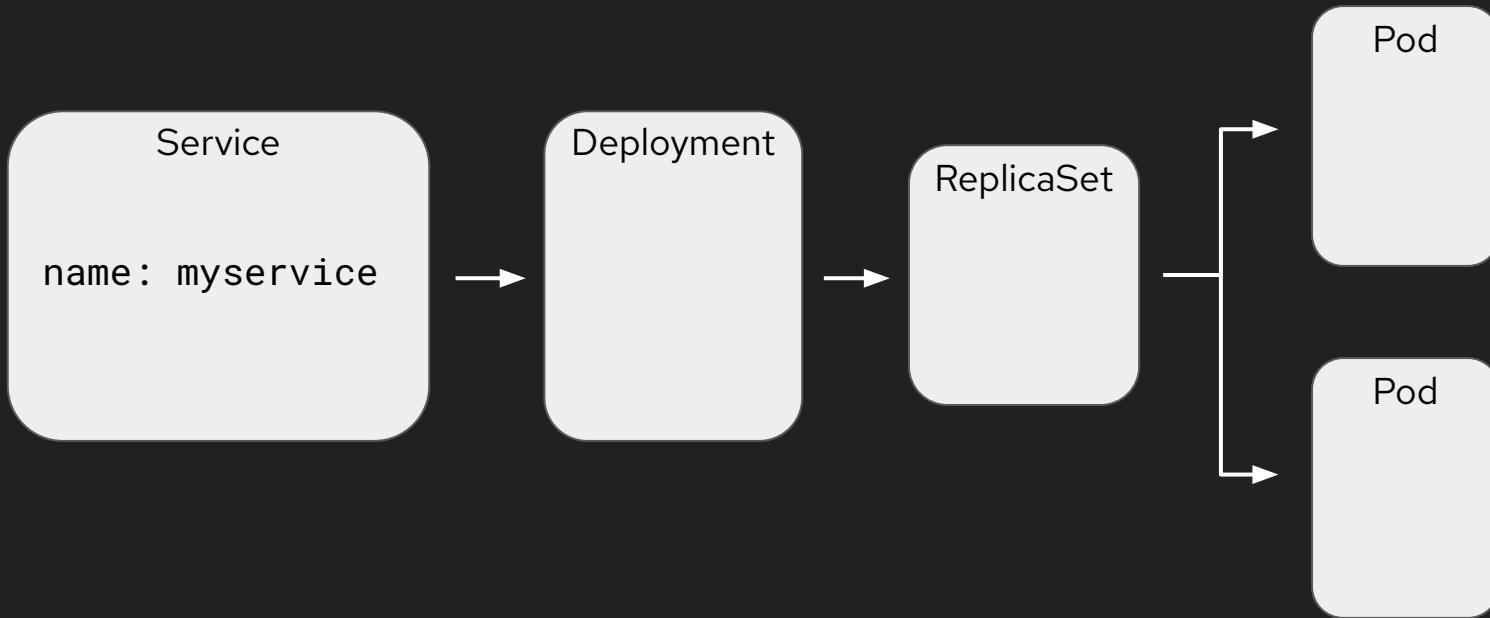
<https://kubernetes.io/docs/tutorials/kubernetes-basics/update/update-intro/>

Service

```
apiVersion: v1
kind: Service
metadata:
  name: myboot
spec:
  ports:
  - name: http
    port: 8080
  selector:
    app: myboot
  type: LoadBalancer
```

Service

Service



kubectl



kubectl commands

<https://kubernetes.io/docs/user-guide/kubectl/>

<https://kubernetes.io/docs/reference/kubectl/cheatsheet>

kubectl get namespaces

kubectl get pods -n mynamespace

kubectl run myboot --image=quay.io/burrsutter/myboot:v1 --port=8080

kubectl logs myboot-75f94d76fc-pl2l7

kubectl expose deployment --port=8080 myboot --type=LoadBalancer

kubectl scale deployment myboot --replicas=3

kubectl set image deployment myboot myboot=quay.io/burrsutter/myboot:v2

Logs



logs

```
System.out.println("Java, where am I?");  
Or console.log("Node logs");
```

```
kubectl get pods  
kubectl logs microspringboot1-2-nz8f8  
kubectl logs microspringboot1-2-nz8f8 -p # last failed pod
```

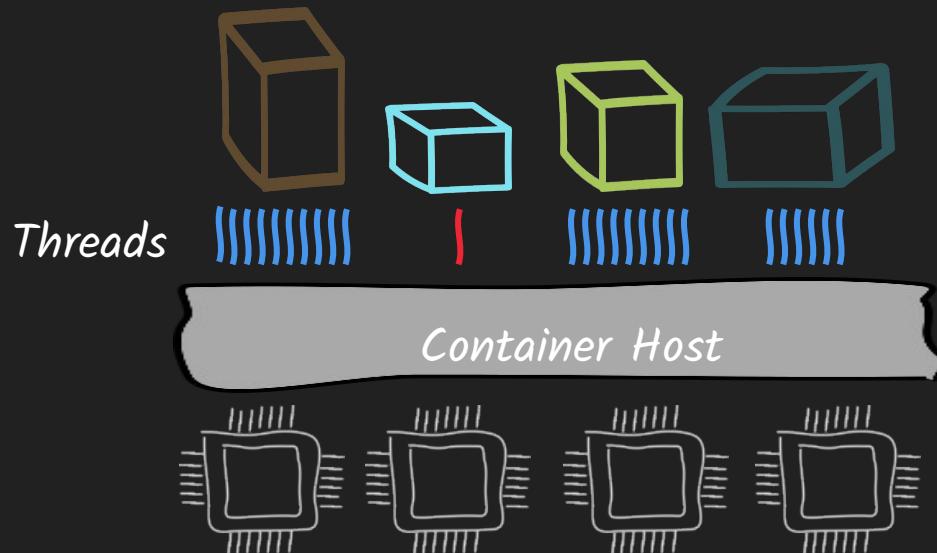
```
OR stern (brew install stern)  
https://github.com/wercker/stern  
stern myboot
```

```
OR kail (https://github.com/boz/kail)  
kail -d myboot
```

Resource Limits

The hidden truth of containers

Containers are about **sharing** and **deployment density**



- | costs memory
- ||||| cost lots of memory
- ||||| cost CPU cycles
- ||||| + quotas = BOOM

Resource Requests/Limits

```
resources:  
  requests:  
    memory: "300Mi"  
    cpu: "250m" # 1/4 core  
  
  limits:  
    memory: "400Mi"  
    cpu: "1000m" # 1 core
```

Liveness & Readiness



“The kubelet uses liveness probes to know when to restart a container.”

“The kubelet uses readiness probes to know when a container is ready to start accepting traffic.”

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-liveness-readiness-startup-probes/>

Example

```
kubectl create -f Deployment.yml
```

```
resources:
  requests:
    memory: "300Mi"
    cpu: "250m" # 1/4 core
  limits:
    memory: "400Mi"
    cpu: "1000m" # 1 core
livenessProbe:
  httpGet:
    port: 8080
    path: /
    initialDelaySeconds: 10
    periodSeconds: 5
    timeoutSeconds: 2
readinessProbe:
  httpGet:
    path: /health
    port: 8080
    initialDelaySeconds: 10
    periodSeconds: 3
```

Environment & ConfigMap



env vars & configmaps

An app's config is everything that is likely to vary between deploys (staging, production, developer environments, etc). [12 Factor Apps](#)

```
kubectl set env deployment/myboot DBCCONN="jdbc:sqlserver://45.91.12.123:1443;user=MyUserName;password=*****;"  
  
kubectl create cm my-config --from-env-file=config/some.properties
```



ConfigMaps

```
kubectl create cm my-config --from-env-file=config/some.properties
```

some.properties

```
GREETING=jambo
LOVE=Amour
```

Partial deployment.yml

```
spec:
  containers:
    - name: myboot
      image: 9stepsawesome/myboot:v1
      ports:
        - containerPort: 8080
      envFrom:
        - configMapRef:
            name: my-config
```

Partial java

```
@Autowired
```

```
private Environment environment;
```

```
String greeting =
environment.getProperty("GREETING","Default");
```

```
String love =
environment.getProperty("LOVE","Default");
```

Secrets



Red Hat
Developer



DevNation



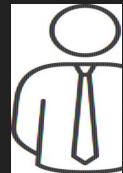
Cambridge
Dictionary

Secret: a piece of information that is only known by one person or a few people and should not be told to others.

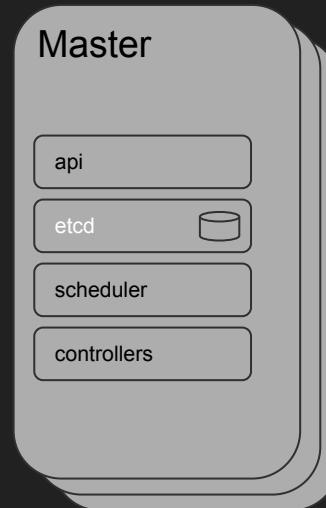
Kubernetes Cluster - Nodes



Dev

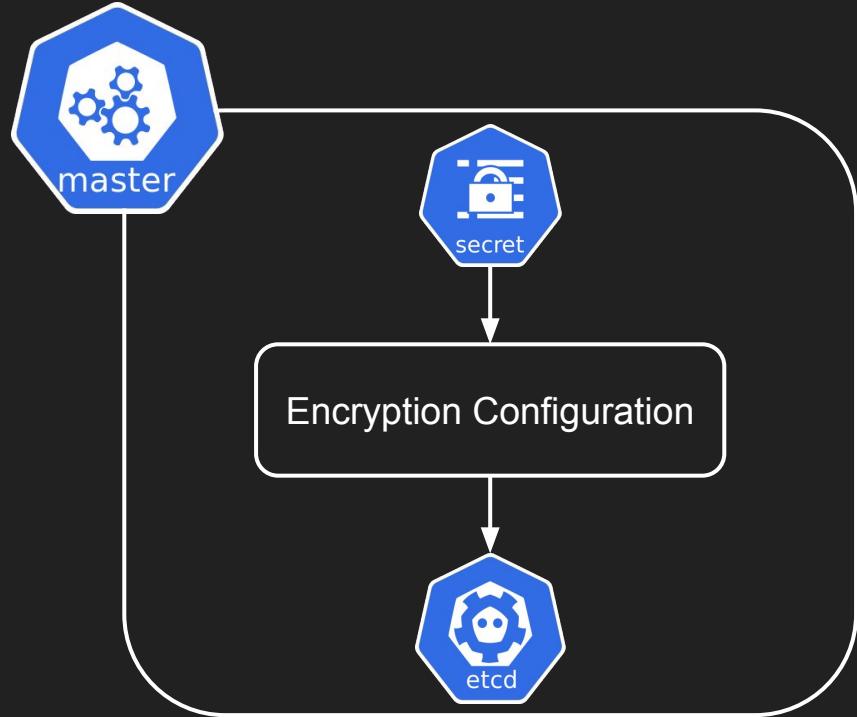


Ops

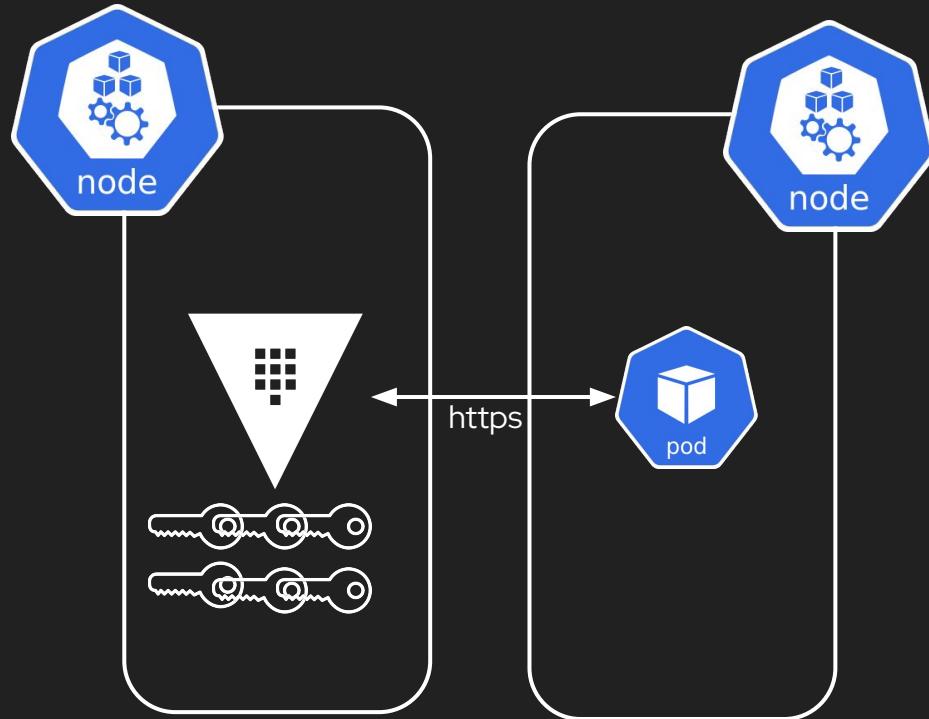




- Kubernetes object that contains a small amount of sensitive data.
- Injected as volume or environment variable.



Vaults provide full encryption



Volumes



The data in volumes is preserved across container restarts.
Lifecycle tied to a pod.



Networked storage. Data is persisted beyond the lifetime of a pod.

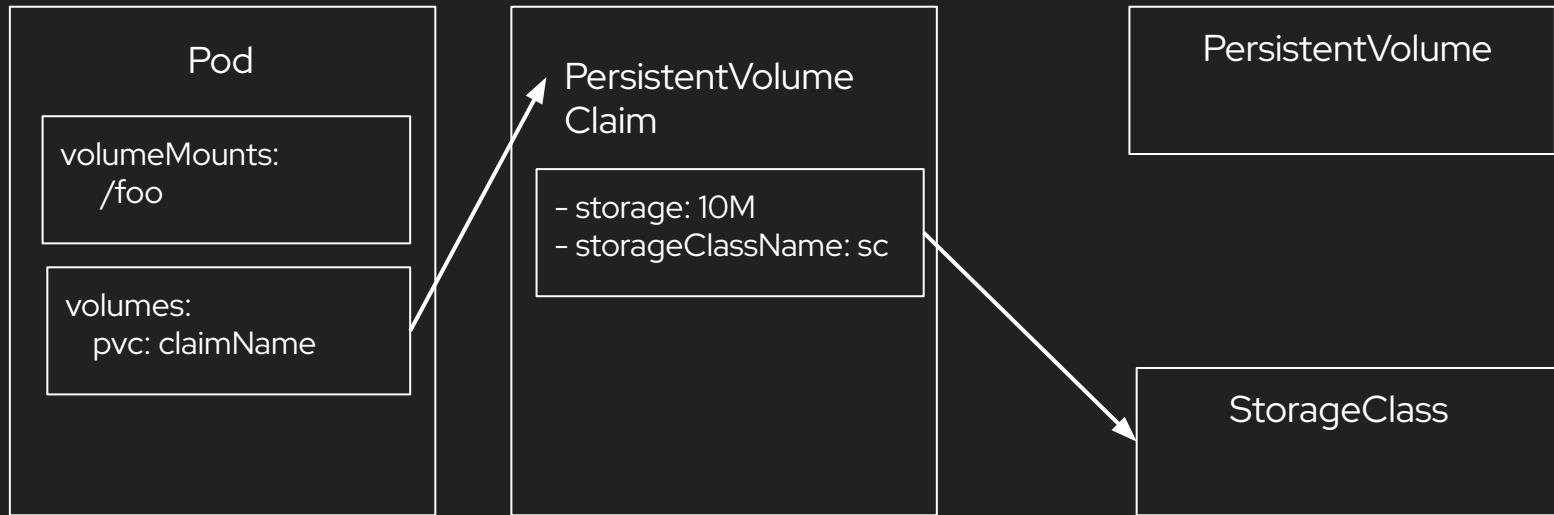


Storage needs to be claimed.



Storage Class for dynamic provisioning of storage.

<https://kubernetes.io/docs/concepts/storage/volumes/>



O'REILLY®

Compliments of

Red Hat
Developer

Introducing Istio Service Mesh for Microservices

Build and Deploy Resilient, Fault-Tolerant Cloud Native Applications



Burr Sutter & Christian Posta

[Download](#)

<https://dn.dev/istiobook>



Red Hat
Developer



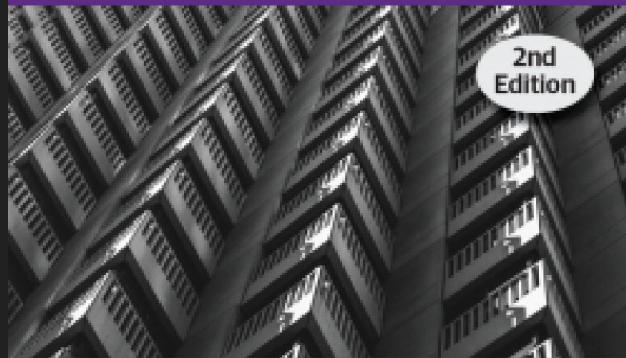
DevNation

O'REILLY®

Compliments of
Red Hat
Developer

Microservices for Java Developers

A Hands-On Introduction
to Frameworks & Containers



Rafael Benevides & Christian Posta

[Download](#)

<https://dn.dev/javamsbook>

 Red Hat
Developer |  DevNation

O'REILLY®

Kubernetes Patterns

Reusable Elements for Designing
Cloud-Native Applications



Bilgin Ibryam &
Roland Huß

 Red Hat
Developer |  DevNation

[Download](#)

<https://dn.dev/k8s-patterns>

O'REILLY®

Knative Cookbook

Building Effective Serverless Applications
with Kubernetes and OpenShift



Burr Sutter &
Kamesh Sampath



Red Hat
Developer

DevNation

[Download](#)

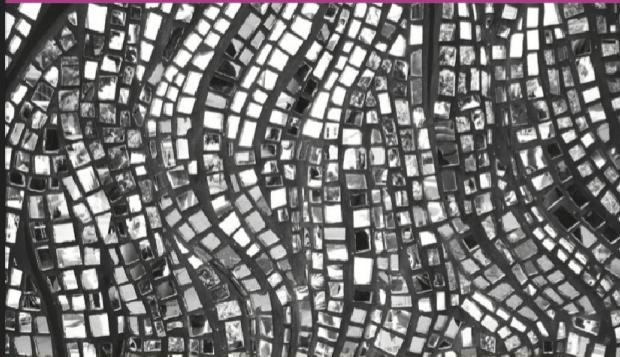
dn.dev/knative-cookbook

dn.dev/kubemaster1

O'REILLY®

Migrating to Microservice Databases

From Relational Monolith
to Distributed Data



Edson Yanaga

Compliments of
Red Hat
Developer

 Red Hat
Developer

 DevNation

[Download](#)

bit.ly/mono2microdb